

Homework #2

Josh Cohen Questions: 4.1,4.2,4.5, 5.2,5.9

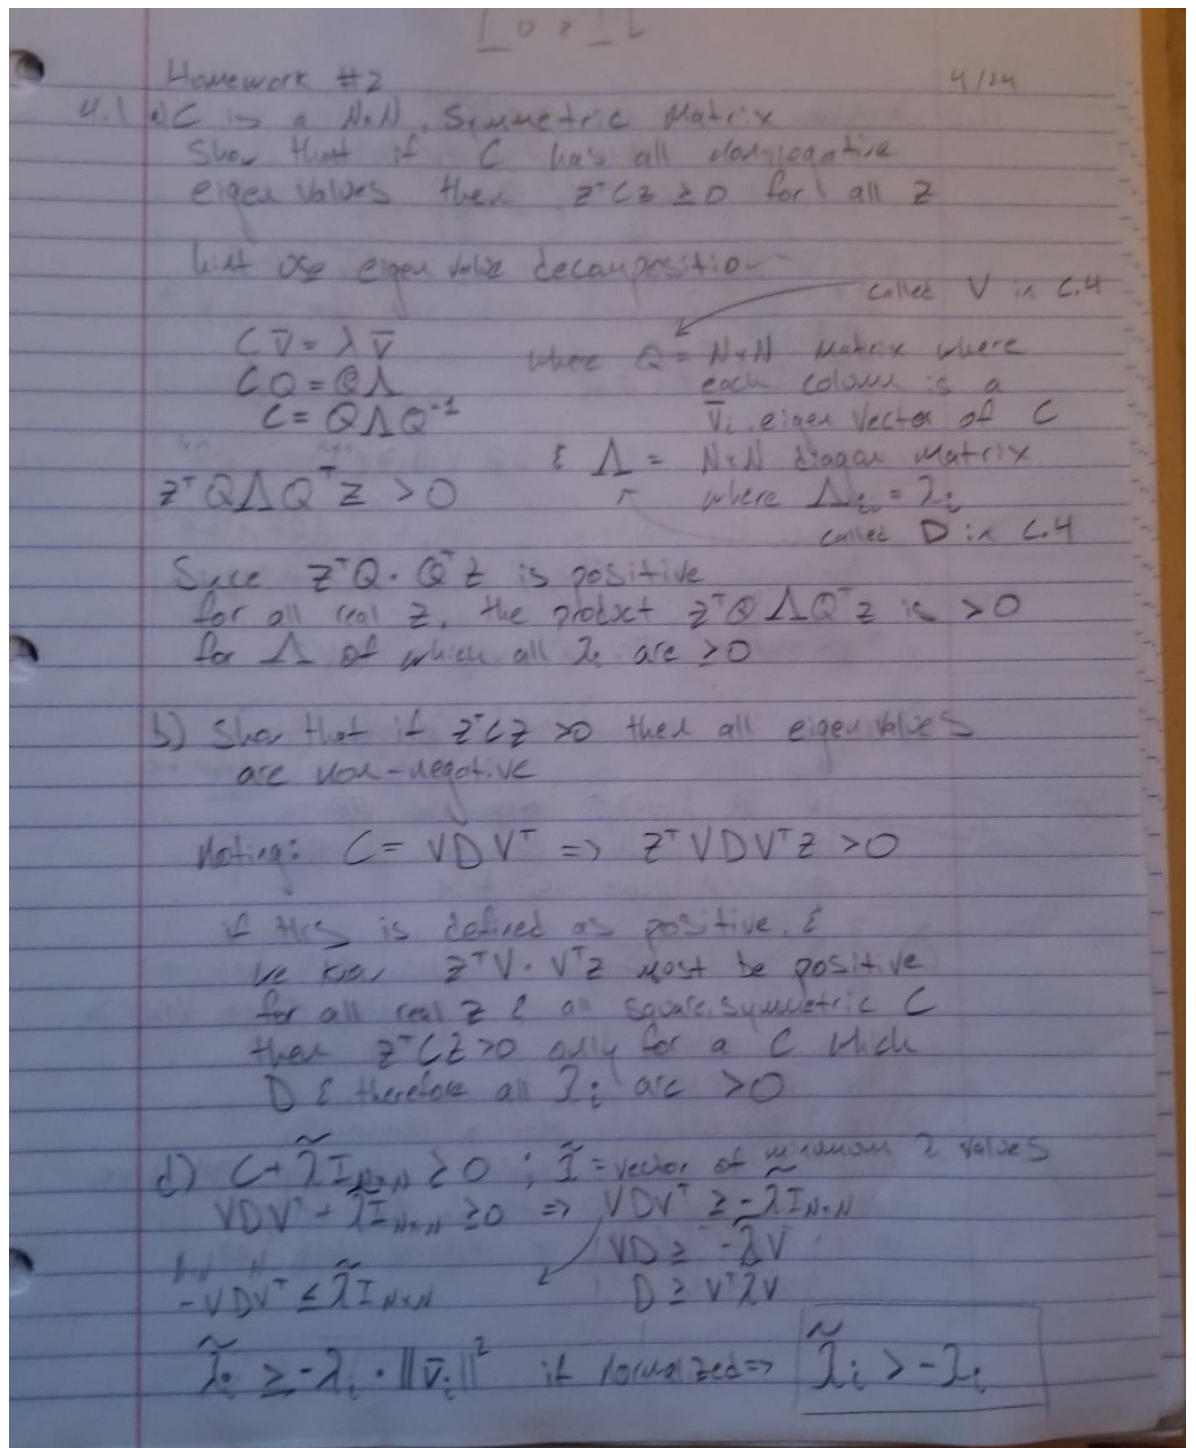
```
In [825]: import numpy as np
import sympy as sym
%matplotlib notebook
from matplotlib import pyplot as plt
import pandas as pd
from IPython.display import Image
```

Problem 4.1

see problems a), b), & d). below

In [828]: Image("4_1.jpg")

Out[828]:



c).

```
In [187]: a = sym.MatrixSymbol('a', 1, 1)
b = sym.MatrixSymbol('b', 2, 1)
C = sym.MatrixSymbol('C', 2, 2)
# w1, w2 = sym.symbols('w_1 w_2')
w = sym.MatrixSymbol('w', 2, 1)
g = a + b.T * w + w.T * C * w
grad_g = g.diff(w)
grad_2_g = grad_g.diff(w)
```

```
In [196]: print("with a g(w) of:")
display(g)
print("gradient g'(w) of:")
display(grad_g)
print("g''(w) of:")
display(sym.factor(grad_2_g.subs({C.T: C})))
```

with a g(w) of:

$$a + b^T w + w^T C w$$

gradient g'(w) of:

$$Cw + b + C^T w$$

g''(w) of:

$$2C$$

```
In [242]: C_act = sym.Matrix([[1,1],[1,1]])
C_act = sym.Matrix(sym.factor(grad_2_g.subs({C.T: C})).subs({C:C_act
}))
print("therefore the second order derivative of our function is:")
display(C_act)
print("with eigen values of:")
eig_vals = C_act.eigenvals(multiple = True)
print(eig_vals)
print("since the eigen values of C are >= to 0 we can verify that the
quadaratic function is convex")
```

therefore the second order derivative of our function is:

$$\begin{bmatrix} 2 & 2 \\ 2 & 2 \end{bmatrix}$$

with eigen values of:

$$[0, 4]$$

since the eigen values of C are >= to 0 we can verify that the quadaratic function is convex

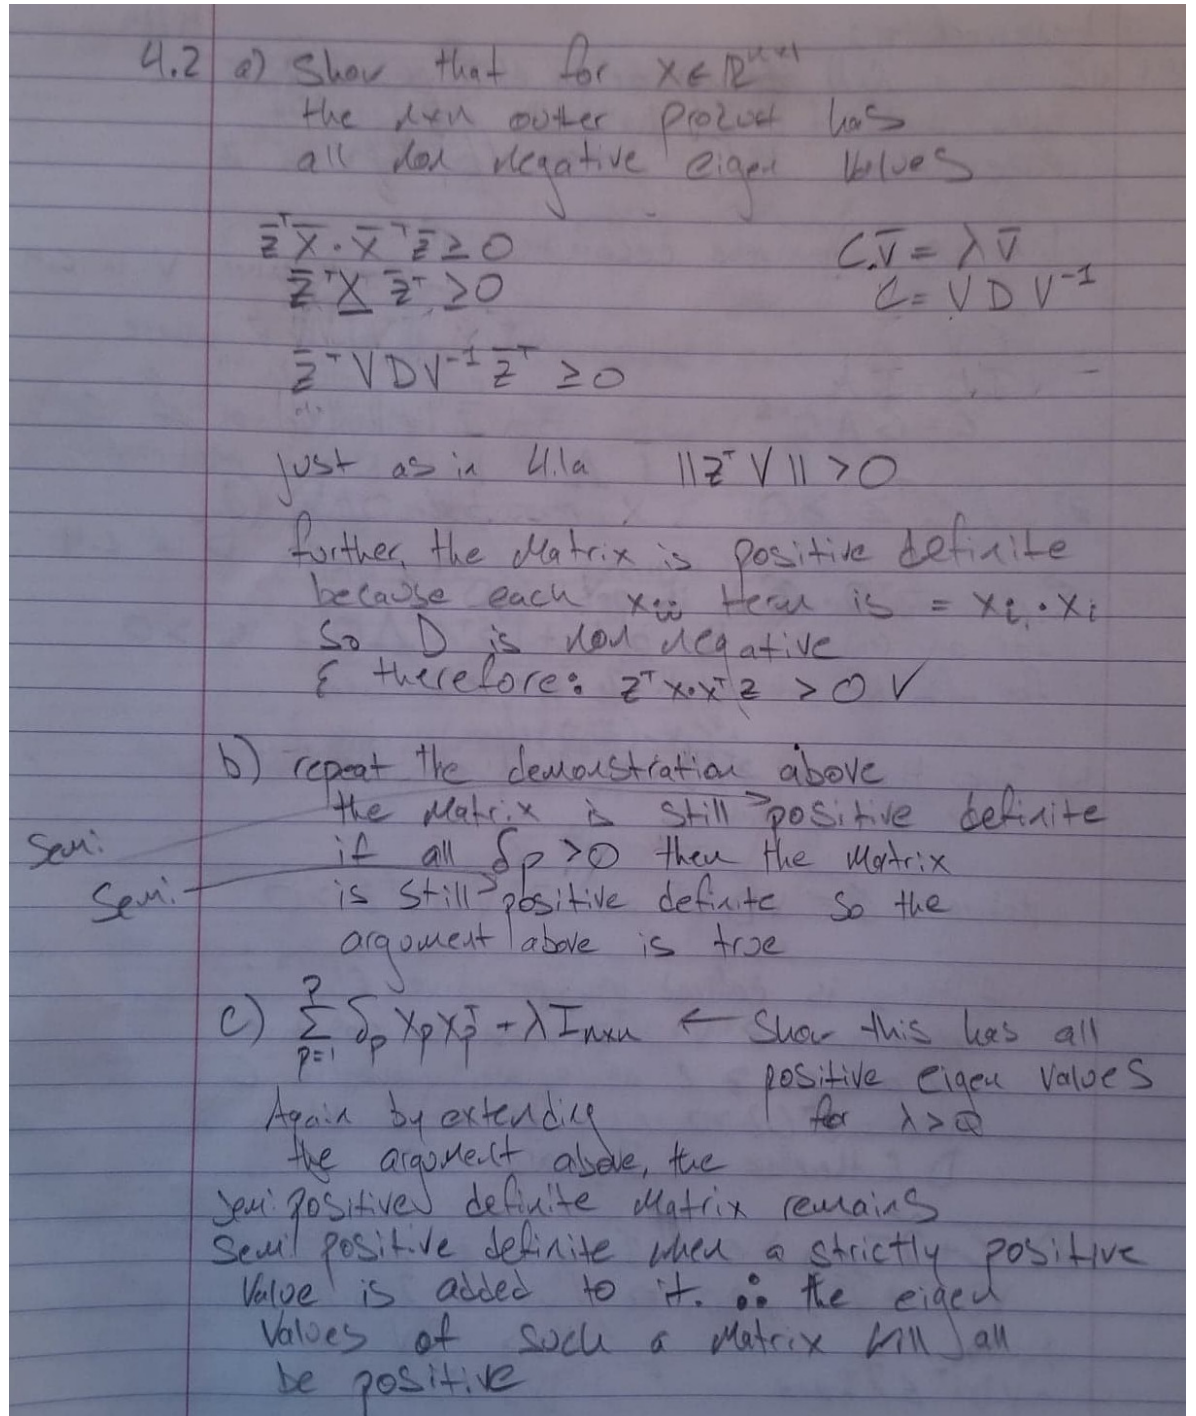
d).

```
In [ ]: C = sym.MatrixSymbol()
```

Problem 4.2

In [829]: Image("4_2.jpg")

Out[829]:



Problem 4.5

a). use 1st order optimality condition to determine the unique stationary point of $g(w) = \log(1 + e^{w^T w})$ where w is 2D

```
In [245]: w1, w2 = sym.symbols('w_1 w_2')
w = sym.Matrix([w1, w2])
g = sym.log(1 + sym.exp(w.T * w)[0])
grad_g = sym.Matrix([g]).jacobian(w)
ans = sym.solve(grad_g, w)
display(ans)
```

$$[(0, 0)]$$

```
In [246]: display(g)
display(grad_g)
```

$$\log(e^{w_1^2 + w_2^2} + 1)$$

$$\begin{bmatrix} \frac{2w_1 e^{w_1^2 + w_2^2}}{e^{w_1^2 + w_2^2} + 1} & \frac{2w_2 e^{w_1^2 + w_2^2}}{e^{w_1^2 + w_2^2} + 1} \\ \frac{2w_1 e^{w_1^2 + w_2^2}}{e^{w_1^2 + w_2^2} + 1} & \frac{2w_2 e^{w_1^2 + w_2^2}}{e^{w_1^2 + w_2^2} + 1} \end{bmatrix}$$

Therefore it can be seen the first order optimality condition is satisfied at $w = (0, 0)$

b). use second order definition of convexity to verify that $g(w)$ is convex and therefore the stationary point is the global minimum

```
In [247]: grad_2_g = grad_g.jacobian(w)
P, D = grad_2_g.diagonalize()
```

```
In [248]: display(D)
```

$$\begin{bmatrix} \frac{2e^{w_1^2 + w_2^2}}{e^{w_1^2 + w_2^2} + 1} & 0 \\ 0 & \frac{2(2w_1^2 + 2w_2^2 + e^{w_1^2 + w_2^2} + 1)e^{w_1^2 + w_2^2}}{2e^{w_1^2 + w_2^2} + e^{2w_1^2 + 2w_2^2} + 1} \end{bmatrix}$$

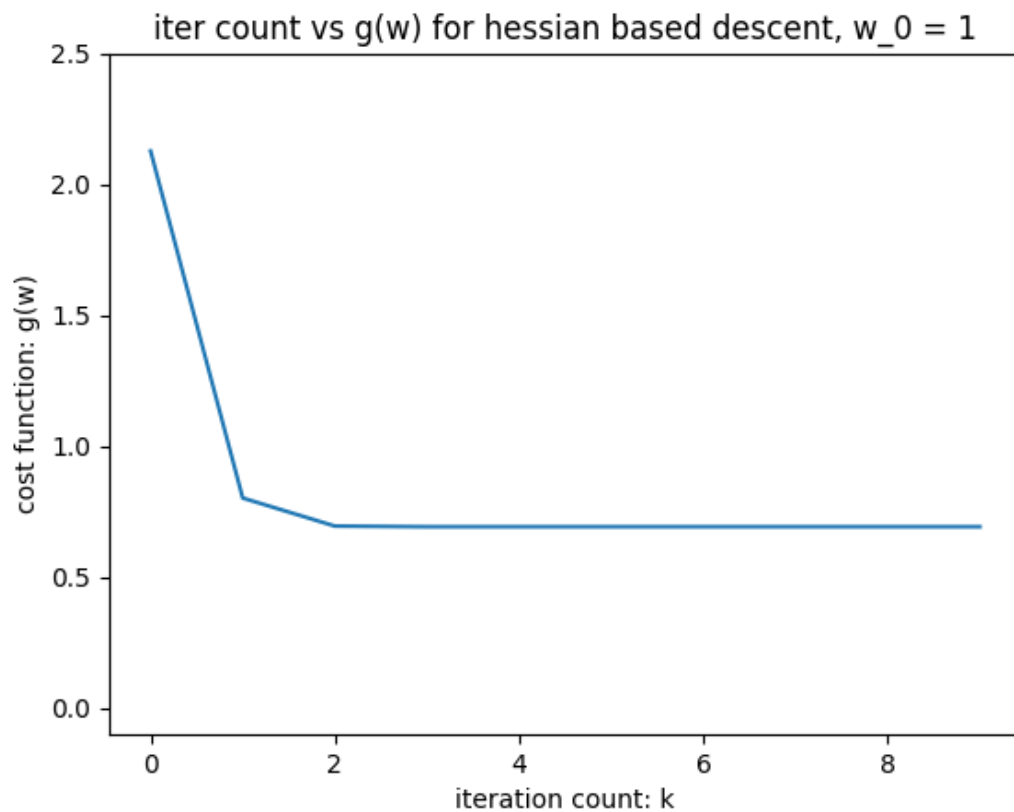
Therefore since e^x can never be < 0 , from the work done in 4.1 and 4.2 we can show that the $N \times N$ symmetric Hessian matrix will never be < 0 , therefore the function is convex everywhere and the minima we found in part **a).** is the global minimum

c). Perform Newton's method to find the minimum of $g(w)$ determined in part **a).**, initialize $w_0 = \mathbf{1}_{N \times 1}$ plot the cost function over 10 iterations

```
In [249]: cost_func = sym.lambdify([w], g)
grad = sym.lambdify([w], grad_g)
grad_2 = sym.lambdify([w], grad_2_g)
w_0_c = np.ones(2)
```


```
In [250]: #define function that we will use for gradient descent
def grad_desc(w_last):
    num_its = 10
    cost_vals = np.zeros(num_its)
    grad_g_last = grad(w_last)
    grad_2_g_last = grad_2(w_last)
    cost_vals[0] = cost_func(w_last)
    #     print(cost_func(w_last))
    #     print(grad_2_g_last.shape)
    #     print(grad_g_last.T.shape)
    for x in range(1, num_its):
        w_next = w_last - np.matmul(np.linalg.inv(grad_2_g_last), grad_g_last.T)[0]
        cost_vals[x] = cost_func(w_next)
        w_last = w_next
        grad_g_last = grad(w_last)
        grad_2_g_last = grad_2(w_last)
    return cost_vals
```

```
In [255]: cost_vals = grad_desc(w_0_c)
plt.figure()
it_vec = np.arange(10)
plt.plot(it_vec, cost_vals)
plt.xlabel("iteration count: k")
plt.ylabel("cost function: g(w)")
plt.ylim([-0.1,2.5])
plt.title("iter count vs g(w) for hessian based descent, w_0 = 1")
plt.show()
```



d). Run code from **c).** again with initial guess at $w_0 = 4 * 1_{N \times 1}$, explain behavior

```
In [244]: w_0_d = np.ones(2) * 4
cost_vals = grad_desc(w_0_d)
plt.figure()
it_vec = np.arange(10)
plt.plot(it_vec, cost_vals)
plt.xlabel("iteration count: k")
plt.ylabel("cost function: g(w)")
plt.title("iter count vs g(w) for hessian based descent, w_0 = 4")
plt.show()
```



why does this result make sense for our $g(w)$? Since our $g(w)$ is a quadratic it makes sense that our second order update rule converges quickly, the one with the greater initial error will then have (due to its globally convex nature) a steeper initial descent which will allow for it to converge quicker and closer to the actual solution

Problem 5.2

a). fit linear model to dataset: *kleibers_law_data.csv*


```

In [470]: #read in data
kliebers_df = pd.read_csv("kleibers_law_data.csv", header=None, index_col=None)
P = kliebers_df.shape[1]
x_vals = kliebers_df.iloc[0,:]
y_vals = kliebers_df.iloc[1,:]
#define symbolic cost function
b, w, xp, yp, p = sym.symbols('b, w xp yp p')
w_tilda = sym.Matrix([b, w])
xp_gen = sym.lambdify(x_sym, sym.Matrix([1, x_sym]))
yp_gen = sym.lambdify(y_sym, sym.Matrix([y_sym]))

#to perform gradient descent calculate A and b_vec
Ap = sym.lambdify(xp, (xp_gen(xp) * xp_gen(xp).T))
bp = sym.lambdify((xp, yp), xp_gen(xp) * yp_gen(yp))
A = np.zeros((2,2))
b = np.zeros((2,1))
for (xp,yp) in zip (np.log(x_vals.iloc[1:]), np.log(y_vals.iloc[1:])):
    A = A + Ap(xp)
    b = b + bp(xp,yp)

```

```

In [471]: display(A)
display(b)

array([[ 1497.          , -2644.10450903],
       [-2644.10450903, 14575.33617811]])

array([[ 8474.33295994],
       [-8512.73814967]])

```

```

In [472]: A_inv = np.linalg.inv(A)
w_star = np.matmul(A_inv,b)
display(w_star)

array([[6.8119607 ],
       [0.65170352]])

```

```

In [473]: w_star[1][0]

```

```

Out[473]: 0.6517035173670688


```

```
In [477]: plt.figure()
ax = plt.gca()
plt.ylabel("log of metabolic rate")
plt.xlabel("log of mass")
plt.scatter(np.log(x_vals), np.log(y_vals))

#plot our fitted line
log_xs = np.log(x_vals)
log_ys = w_star[1][0] * log_xxs + w_star[0][0]
plt.plot(log_xs, log_ys, 'r')

plt.legend(["data set", "best fit"])

plt.show()
```



b)... right out nonlinear relationship between mass x and metabolic rate y

```
In [480]: slope = w_star[1][0]
intercept = w_star[0][0]
print("best fit parameters for:\n\rm: {} \r\nb: {}".format(slope, intercept))
```

```
best fit parameters for:
m: 0.6517035173670688
b: 6.811960700573648
```

writing this out as a non-linearized equation:

$$y_p = e^{w_0 + \log x_p w_1}$$

or

$$y_p = e^{w_0} x_p^{w_1}$$

c). plug in a mass of 10kg to find out how many calories a 10kg animal needs daily

```
In [513]: animal_mass = 10 #in kg
daily_cals = np.exp(intercept) * animal_mass**slope * 1000/4.18 #in cal/day
```

```
In [521]: print("There for the 10kg animal needs: {:.2f} calories daily;\nHowever, in our 'calorie' which is infact a kilocalorie this number is: {:.2f} 'calories' a day".format(daily_cals, daily_cals/1000))
```

There for the 10kg animal needs: 974819.43 calories daily;
However, in our 'calorie' which is infact a kilocalorie this number is: 974.82 'calories' a day

Problem 5.9

Standardize and normalize quality metrics, verify example 5.5 and 5.6

```
In [526]: #read in data
boston_housing_df = pd.read_csv("boston_housing.csv", header=None, index_col=None).dropna(axis=1)
auto_data_df = pd.read_csv("auto_data.csv", header=None, index_col=None).dropna(axis=1)
```

```
In [609]: #standardize and normalize
bh_mean = boston_housing_df.mean(axis = 1)
ad_mean = auto_data_df.mean(axis=1)
bh_std = boston_housing_df.std(axis=1)
ad_std = auto_data_df.std(axis=1)

bh_std_norm = boston_housing_df.sub(bh_mean, axis=0).div(bh_std, axis=0)
ad_std_norm = auto_data_df.sub(ad_mean, axis=0).div(ad_std, axis=0)
```

```

In [740]: def get_best_fit(df):
            inputs = df.iloc[:-1,:]
            output = df.iloc[-1,:]
            N, P = inputs.shape
            b, yp_sym = sym.symbols('b yp')
            w = sym.MatrixSymbol('w', N, 1).as_explicit() #first element is b
            w_tilda = sym.Matrix([b, w])
            xp_sym = sym.MatrixSymbol('x', N,1)
            xp_gen = sym.lambdify([xp_sym],sym.Matrix([sym.Matrix([1]), xp_sym]))
            yp_gen = sym.lambdify(yp_sym, sym.Matrix([yp_sym]))
            #to perform gradient descent calculate A and b_vec
            # Ap = sym.lambdify([xp_sym], (sym.MatMul(xp_gen(xp_sym),xp_gen(xp_sym).T)))
            bp = sym.lambdify([xp_sym], yp_sym), xp_gen(xp_sym) * yp_gen(yp_sym))
            A = np.zeros((N+1,N+1))
            b = np.zeros((N+1,1))
            for i in range(1,P):
                xp = xp_gen(np.expand_dims(inputs.iloc[:,i].to_numpy(),1))
                yp = output.iloc[i]
                A = A + np.matmul(xp,xp.T)
                b = b + xp * yp

            A_inv = np.linalg.inv(A)
            w_star = np.matmul(A_inv, b)
            x_real = np.concatenate((np.ones((P,1)).T, inputs.to_numpy()))
            y_real = np.expand_dims(output.to_numpy(),1)
            return (w_star, x_real, y_real)

```

```

In [795]: def get_errors(w_star, x_real, y_real):
            y_theoro = np.matmul(x_real.T, w_star)
            y_diff = np.abs(np.subtract(y_theoro,y_real))
            y_diff_sq = np.multiply(y_diff,y_diff)
            RMSE = np.sqrt(np.mean(y_diff_sq))
            MAD = np.mean(y_diff)

            return (RMSE, MAD)

```

```

In [797]: w_star, x_real, y_real = get_best_fit(bh_std_norm)
            bh_RMSE, bh_MAD = get_errors(w_star, x_real, y_real)
            w_star, x_real, y_real = get_best_fit(ad_std_norm)
            ad_RMSE, ad_MAD = get_errors(w_star, x_real, y_real)

```

```
In [824]: print("Therefore we have the following errors for our best fits:\r\n"
)
print("Boston Housing:\r\nRMSE: {:.5f}\r\nMAD: {:.5f}".format(bh_RMSE
, bh_MAD))
print("Actual RMSE/MAD: {:.5f}\r\nTarget RMSE/MAD: {:.5f}".format(bh_
RMSE/bh_MAD, 4500/3000 ))
print("Actual RMSE/Theoretical RMSE: {:.5f}".format(4500/bh_RMSE))
print("Actual MAD/Theoretical MAD: {:.5f}".format(3000/bh_MAD))
print("")
print("Auto Data:\r\nRMSE: {:.5f}\r\nMAD: {:.5f}".format(ad_RMSE, ad_
MAD))
print("Actual RMSE/MAD: {:.5f}\r\nTarget RMSE/MAD: {:.5f}".format(ad_
RMSE/ad_MAD, 3.3/2.5))
print("Actual RMSE/Theoretical RMSE: {:.5f}".format(3.3/ad_RMSE))
print("Actual MAD/Theoretical MAD: {:.5f}".format(2.5/ad_MAD))
```

Therefore we have the following errors for our best fits:

Boston Housing:

RMSE: 0.50882

MAD: 0.35624

Actual RMSE/MAD: 1.42832

Target RMSE/MAD: 1.50000

Actual RMSE/Theoretical RMSE: 8844.05012

Actual MAD/Theoretical MAD: 8421.39522

Auto Data:

RMSE: 0.42199

MAD: 0.32036

Actual RMSE/MAD: 1.31724

Target RMSE/MAD: 1.32000

Actual RMSE/Theoretical RMSE: 7.82009

Actual MAD/Theoretical MAD: 7.80373

note

the results above should be satisfactory evidence that the best fits were found

In []: